

alexandria.

investment memo

Most AI investments cannot defend their terminal value.

Alexandria's compounds with every model release — especially post-singularity.

The cognitive layer of AI — where human minds remain machine-readable across the transition.

the investment.

Most AI products are transition assets — value decays into the model with each release. Alexandria's asset is a **private, accumulating, machine-readable record of how a specific person thinks** — sovereign to the user, unsynthesizable from outside, *worth more from every better model*. The strategic prize is the canonical write path for those records: same shape as WordPress, GitHub, Hugging Face. Winner-takes-most. Coding throughput is not the constraint. **Benjamin has personal runway through the end of 2026; distribution is the current constraint, and capital is needed before runway ends.** Engineering remains founder-led through coding agents, with more model usage or a bounded specialist engagement only if scale earns it.

In April 2026 Andrej Karpathy publicly described this architecture — raw data compiled into structured markdown, incrementally enhanced, queries that compound back into the knowledge base — and said "there is room here for an incredible new product." Alexandria had already shipped it.

the asset.

A structured record of who each user is — whatever they choose to put in it: worldview, values, mental models, taste, personal thoughts, decisions, blind spots, anything. **The product is the framework; the practice is theirs** (daily Socratic sessions, journals, decision logs, any system works). Local files the user owns. Portable across every AI provider. *The record cannot be back-filled: inbox-scraping reaches what the user does, never what they mean.* **The compounding is the years.**

The product reached v3; real-user hardening begins now. One paste. The tool is free and yours; membership is \$30/month after a 30-day trial, or free while three referred friends remain active. The live questions are who starts, stays and pays, and what breaks when they do.

the moat.

Alexandria is sovereign by design — any individual user can take their files and walk at zero cost. **The defensibility is one layer up: coordination.** *Alexandria only works as the read **and** write path for personal thought — writers chase readers, readers chase writers, both terminate on the same place.* Once the network gathers there, displacing Alexandria requires getting everyone to switch at once. That doesn't happen — same as HTTP, SMTP, GitHub, Hugging Face. **The scarce asset in a world of cheap intelligence is not capability; it is coordination.**

1. Convention. *Be the canonical place where people host their cognitive records. Same shape as build-it-on-WordPress, push-it-to-GitHub, host-it-on-Hugging-Face.* **Apple can build the tool; Apple cannot build the cathedral.** *Winner-takes-most. 12–36 month window — move as if 12.*

2. Network. *Your circle on Alexandria means your AI can read their frames when planning with or about them.* The friend's-mind primitive turns recruitment into self-interest. *A clone starts with no population.*

3. Integration. *Every AI tool reading from the user's local Alexandria directory is another terminal point.* Claude Code, Cursor, Codex, Grok CLI today; every personalised app tomorrow. **The user can leave; the ecosystem points back.**

the business.

Membership. *The tool and methodology are free; membership is \$30/month after a 30-day trial, or free while three referred friends remain active.* Stands alone. Break-even at 4 paying members.

Platform. *Users grant other services permissioned access to their Alexandria record — opt-in, revocable.* AI labs pay for alignment and personalisation data. Consumer apps pay to know users who want to be known (meditation, therapy, dating, productivity). **The user monetises their own frame; Alexandria captures a clearing fee.** *Independent of membership count.*

If	Then
~100K paying members × \$50/m settled	~\$60M ARR
Platform / integration revenue	~\$40M+ ARR
Combined at standard 10× SaaS multiple	\$1B+ valuation

the founder.

Benjamin a. Mowinckel. 26, Norwegian-American. Eton, Princeton, Norwegian military, brief stint at Verdane. Currently in San Francisco. **Solo, running the entire stack on AI agents at \$100/month opex, full-time since March 2026.** He reached v3 and began external manual onboarding on 24 August 2026.

	\$/m
Company opex	\$100
Founder bare minimum	\$500
Founder safe monthly	\$5,000

Benjamin has personal runway through the end of 2026. *Investor routes fire after two clean outside-user transfers; the round then funds the chosen 2027 runway and proven distribution.*

"I could be wrong. I doubt it. The company costs \$100 a month. If I am wrong, the world loses nothing. If I am right and nobody builds this, the world loses something it did not know it needed. Either way, I am building it."

the ask.

- **A small round before the end of 2026.** Set the exact amount from the desired 2027 runway, first-user evidence, and any distribution motion that earns funding.
- **Engineering remains founder-led.** More model usage comes first; use a paid one-off or part-time specialist only for a bounded problem that survives it.

The company needs relatively little money, and the founder is starting before runway pressure.

*The firm whose name attaches goes on the page about how humans preserved thinking through the AI transition. **America protects the right to think. Alexandria protects the ability.** Set the angels free.*

Benjamin a. Mowinckel · benmowinckel@gmail.com · alexandria-library.com · a.