

alexandria.

investment memo

Own your progress. Change your AI. Choose what to share.

The cognitive layer of AI — where human minds remain machine-readable across the transition.

the investment.

People should be free to change their AI without starting their personal system over. They own the persistent data and intent; models, harnesses and tools can change around it. Benjamin built an Alexandria loop for himself and open-sourced it so anyone can adopt the parts that fit. The person's own AI reads, develops and writes their private record. **That record is the person's asset, never the company's.** The company begins at the optional connection between those independently owned systems: people deliberately publish selected Mirrors, and Alexandria maintains discovery, verified identity and current permissions. They can keep their own website, storage and model. **We help you build something you own. We charge for connecting it to other people who own theirs.** Software and a registry are reproducible; a lone individual cannot supply other people's continuing participation, contributions and permission. Coordination and Alexandrian identity could make the service harder to replace as useful participation grows. **There is no moat today, and better AI does not guarantee one.** Coding throughput is not the constraint. **Benjamin has personal runway through the end of 2026; distribution is the current constraint, and capital is needed before runway ends.** *Engineering remains founder-led through coding agents, with more model usage or a bounded specialist engagement only if scale earns it.*

In April 2026 Andrej Karpathy publicly described this architecture — raw data compiled into structured markdown, incrementally enhanced, queries that compound back into the knowledge base — and said "there is room here for an incredible new product." Alexandria had already shipped it.

the asset.

The person's own AI reads, develops and writes a record of their thinking in files they own. Benjamin open-sourced his own Alexandria loop so people can use whichever parts fit their existing setup. **Capture, Loop, Skill and Mirror recipes are free to copy and keep.** The private map supports personalisation. Each person chooses what becomes a Mirror: selected writing on their own website, or an optional live PLM using their own model. Any AI with authorised access can read the owned files; hidden vendor state is outside that promise. **The paid Connector links those selected Mirrors for useful discovery and permissioned access.** That is cross personalisation: your own AI uses your context with another person's permitted depth.

The product reached v3; real-user hardening begins now. One paste. The tool is free and yours; membership is \$30/month after a 30-day trial, or free while three referred friends remain active. The live questions are who starts, stays and pays, and what breaks when they do.

the moat.

Alexandria is sovereign by design: people keep their private files, independent websites and model accounts when they leave. Free recipes are theirs to keep. **The company has two possible sources of defensibility: coordination and brand-identity. Neither is built yet.** Coordination means a useful place to find particular people and access the depth they choose to share; brand-identity means people value being Alexandrian and continuing the practice. A competitor can reproduce the code and published snapshots. It cannot create other people's fresh contributions or permission merely by copying the software. That gives a participating collective potential value, not guaranteed dominance. People can leave, use several networks, or coordinate elsewhere. **The service must keep making permitted connections materially easier.** A directory of shallow biographies and chat wrappers fails that test. The company operates shared discovery, identity and exact current access; it need not host every page, hold private records or relay independent public model calls. **Build for useful participation, not a captive dataset.**

1. Convention. *Become a useful shared connection underneath independently owned websites.* A hosted profile is an optional starting point. A person can keep their own pages, published files and model; no mandatory company host or private write path is required.

2. Network. *Your own AI can use your context alongside the parts other people choose to share. Useful depth, voluntary contribution and current permissions create the value. A software clone begins without those people’s participation.*

3. Integration. *An existing website can add a public description of selected files, then optionally add the shared Connector or live chat. The private loop works without the company. Independent public reads and own-model calls need not pass through Alexandria.*

the business.

Membership. *The tool and methodology are free; membership is \$30/month after a 30-day trial, or free while three referred friends remain active. Pays for the maintained connection, not ownership of the free loop. Break-even at 4 paying members.*

Platform. *Users grant other services permissioned access to their Alexandria record — opt-in, revocable. AI labs pay for alignment and personalisation data. Consumer apps pay to know users who want to be known (meditation, therapy, dating, productivity). **The user monetises their own frame; Alexandria captures a clearing fee.** Independent of membership count.*

If	Then
~100K paying members × \$50/m settled	~\$60M ARR
Platform / integration revenue	~\$40M+ ARR
Combined at standard 10× SaaS multiple	\$1B+ valuation

the founder.

Benjamin a. Mowinckel. 26, Norwegian-American. Eton, Princeton, Norwegian military, brief stint at Verdane. Currently in San Francisco. **Solo, running the entire stack on AI agents at \$100/month opex, full-time since March 2026.** He reached v3 and began external manual onboarding on 24 August 2026.

	\$/m
Company opex	\$100
Founder bare minimum	\$500
Founder safe monthly	\$5,000

Benjamin has personal runway through the end of 2026. *Investor routes fire after two clean outside-user transfers; the round then funds the chosen 2027 runway and proven distribution.*

"I could be wrong. I doubt it. The company costs \$100 a month. If I am wrong, the world loses nothing. If I am right and nobody builds this, the world loses something it did not know it needed. Either way, I am building it."

the ask.

- **A small round before the end of 2026.** Set the exact amount from the desired 2027 runway, first-user evidence, and any distribution motion that earns funding.
- **Engineering remains founder-led.** More model usage comes first; use a paid one-off or part-time specialist only for a bounded problem that survives it.

The company needs relatively little money, and the founder is starting before runway pressure.

*The firm whose name attaches goes on the page about how humans preserved thinking through the AI transition. **America protects the right to think. Alexandria protects the ability.** Set the angels free.*

Benjamin a. Mowinckel · benmowinckel@gmail.com · alexandria-library.com · a.